

# Indexing

INF 551

Wensheng Wu

# Outline

- Types of indexes
- B+ trees

# Indexes

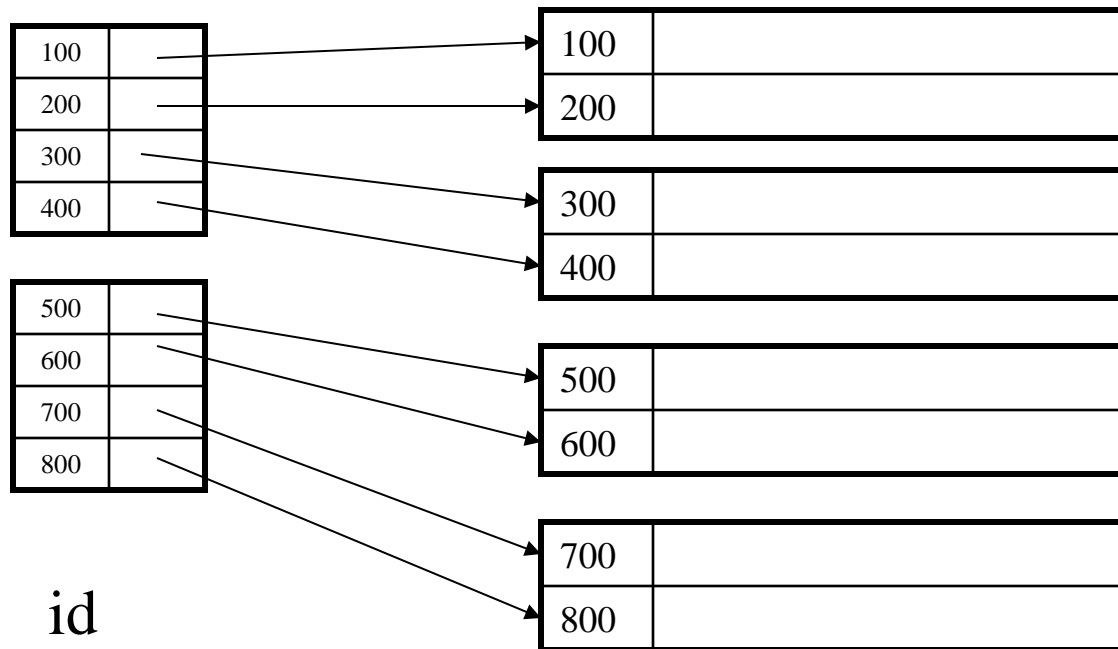
- An *index* is a data structure that speeds up selections on the *search key field(s)*
- *Fields = attributes*
- Search key = any subset of the fields of a relation
  - *Search key* is **not** the same as *key* (minimal set of fields that uniquely identify a record in a relation).
- Entries in an index:  $(k, r)$ , where:
  - $k$  = the key
  - $r$  = the record OR record id OR a list of record ids

# Index Classification

- Clustered/unclustered
  - Clustered = records sorted & stored in the order of search key
  - Unclustered = records are not sorted in key order
- Dense/sparse
  - Dense = each record has an entry in the index
  - Sparse = only some records have
- B+ tree / hash table / ...

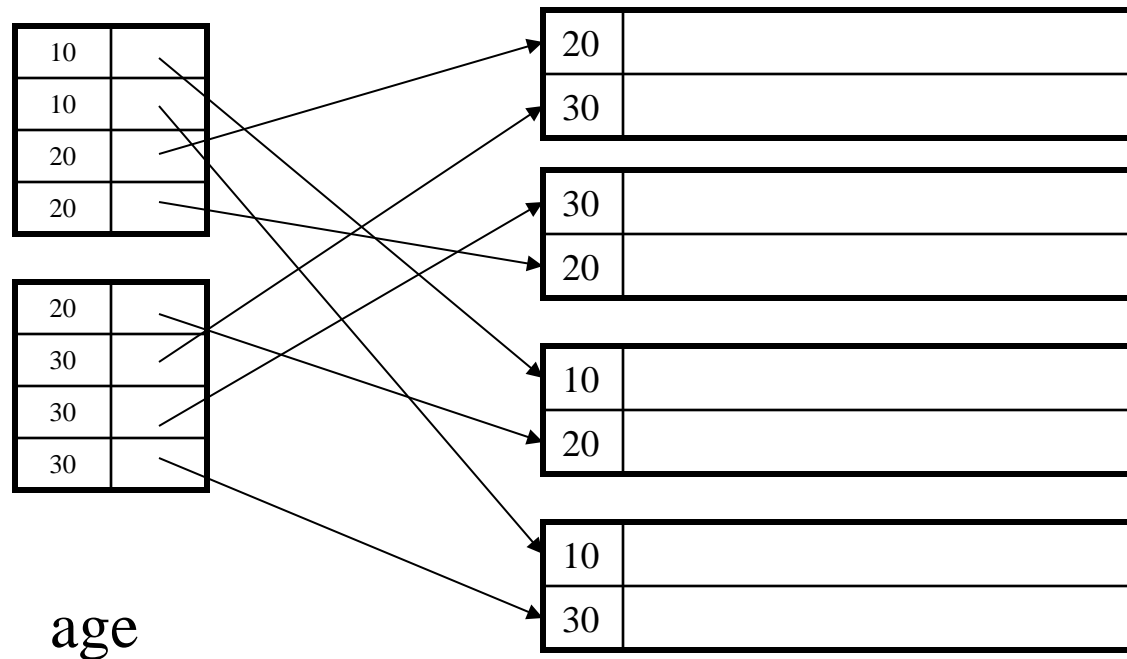
# Clustered Index

- Records are sorted on the index attribute
  - Often for index on primary key
  - E.g., employee(id, name, age, salary)




# Unclustered Indexes

- Often for indexes on attributes other than primary key



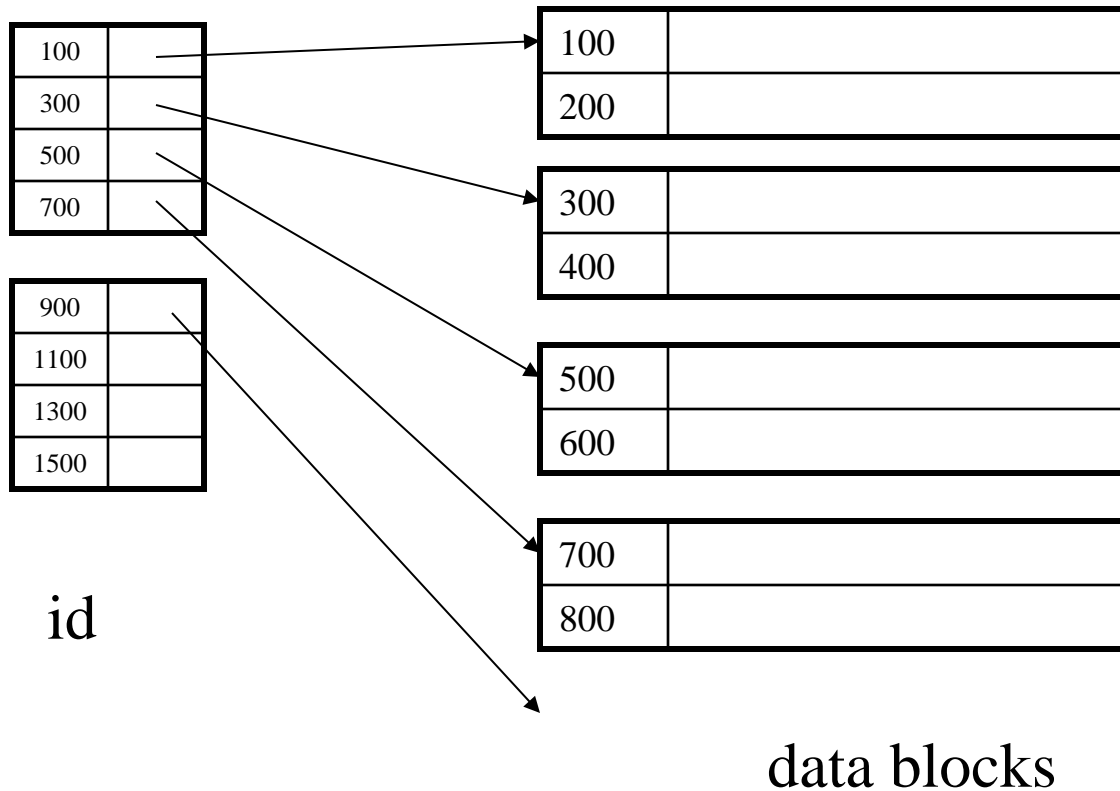
# Dense Index

- Dense index: one key per data record
- See Sections 14.1.2 – 14.1.3 in book [GVW]



# Sparse Index

- Sparse index: one key per data block






# Query Types

- Equality query:  $\langle \text{attribute} \rangle = \langle \text{value} \rangle$ 
  - E.g.,  $\text{age} = 20$ ,  $\text{sal} = 75$
- Range query:  $\langle \text{attribute} \rangle \langle \text{inequality operator} \rangle \langle \text{value} \rangle$ 
  - Inequality operator:  $<$ ,  $>$ ,  $\leq$ ,  $\geq$
  - E.g.,  $\text{age} > 20$  or  $\text{sal} \leq 75$

# Composite Search Keys

- Composite Search Keys*: Search key = a list of fields.


Keys in index  
sorted by  $\langle \text{age}, \text{sal} \rangle$ :  
i.e., first by age; if ties,  
by sal



Keys sorted by  $\langle \text{sal} \rangle$

# Questions

- Which index is useful for queries:
  - $Sal > 75$
  - $Age = 12$  and  $sal > 10$
  - $Age > 12$



# Outline


- Types of indexes
- B+ trees

# B+ Trees


- Search trees
- Idea in B Trees:
  - make 1 node = 1 block
- Idea in B+ Trees:
  - Make leaves into a linked list
  - Efficiently support range queries

# B+ Trees Basics

- Parameter  $d$  = the **degree** (also called order)
- Each node has  $\geq d$  and  $\leq 2d$  keys (except root)




- Each leaf has  $\geq d$  and  $\leq 2d$  keys:



# B+ Tree Example

$d = 2$



# B+ Tree Design

- How large is  $d$  ?
- Example:
  - Key size = 4 bytes
  - Pointer size = 8 bytes
  - Block size = 4096 bytes
- $2d * 4 + (2d+1) * 8 \leq 4096$
- $d = 170 (\sim 170.33)$



# B+ Trees in Practice

- Typical order  $d = 100$ .
- Typical fill-factor (minimum in practice): 66.7% (i.e.,  $2/3$ ) (note minimum fill factor in design: 50%)
  - Minimum # of keys in a node = 133 ( $200 * 2/3$ )
- Capacities (# of records which the index supports):
  - Height 1 (tree with a single root): 133 records
  - Height 2:  $133^2 = 17,689$  records ( $134 * 133$  to be exact)
  - Height 3:  $133^3 = 2,352,637$  records ( $134^2 * 133$ )
  - Height 4:  $133^4 = 312,900,721$  records ( $134^3 * 133$ )

# B+-tree in Practice

- Can often hold top levels in buffer pool:
  - Level 1 = 1 page = 4KB
  - Level 2 = 133 pages = 532KB
  - Level 3 = 17,689 pages = 70,756KB ~ 70MB

# Searching a B+ Tree


- Equality search:
  - Start at the root
  - Proceed down, to the leaf
- Range query:
  - Finding the first leaf
  - Then sequential traversal of leaves

```
Select name  
From people  
Where age = 25
```

```
Select name  
From people  
Where 20 <= age  
and age <= 30
```

# Example


$20 \leq \text{age}$  and  $\text{age} \leq 55$



# Insertion into a B+ Tree

Insert (K, P)

- Find leaf where K belongs, insert
- If no overflow ( $2d$  keys or less), stop
- If overflow ( $2d+1$  keys), split node, insert middle into parent:




# Notes

- Splitting of leaf may lead to splitting of its parent and ancestors
- When splitting a leaf, middle key (e.g., K3) is also kept in the new node on the right
- No need to retain middle key in the split node when splitting an internal node (but need to insert it into parent)
- When root is split, new root has only one key

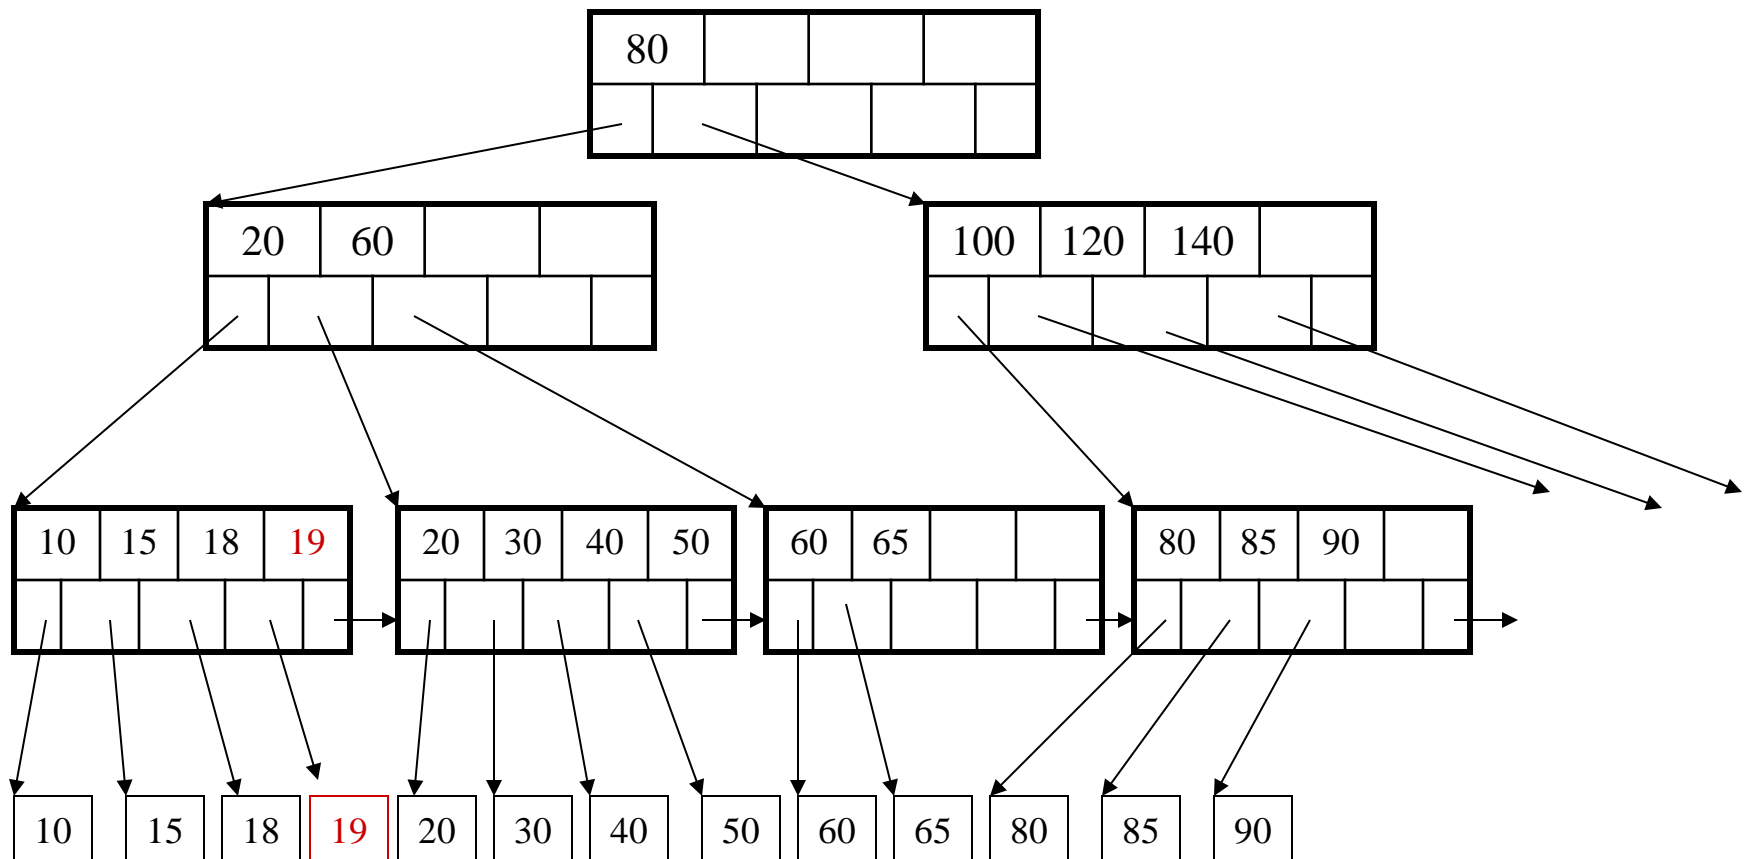
# Insertion into a B+ Tree

Insert K=19



# Insertion into a B+ Tree


After insertion





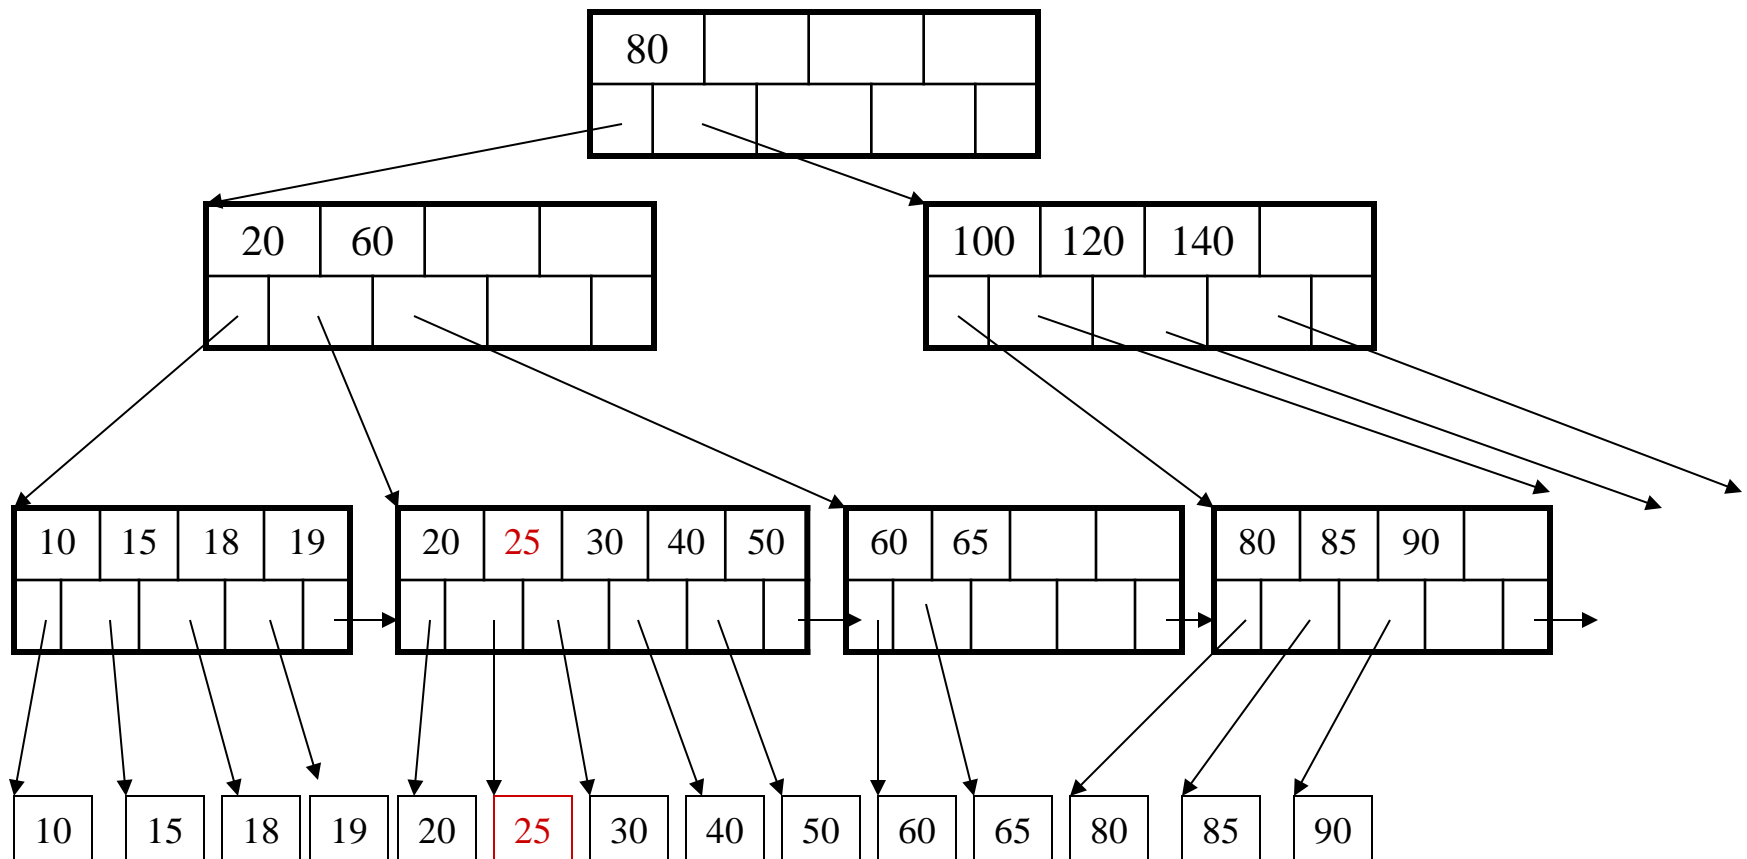
# Insertion into a B+ Tree

Now insert 25



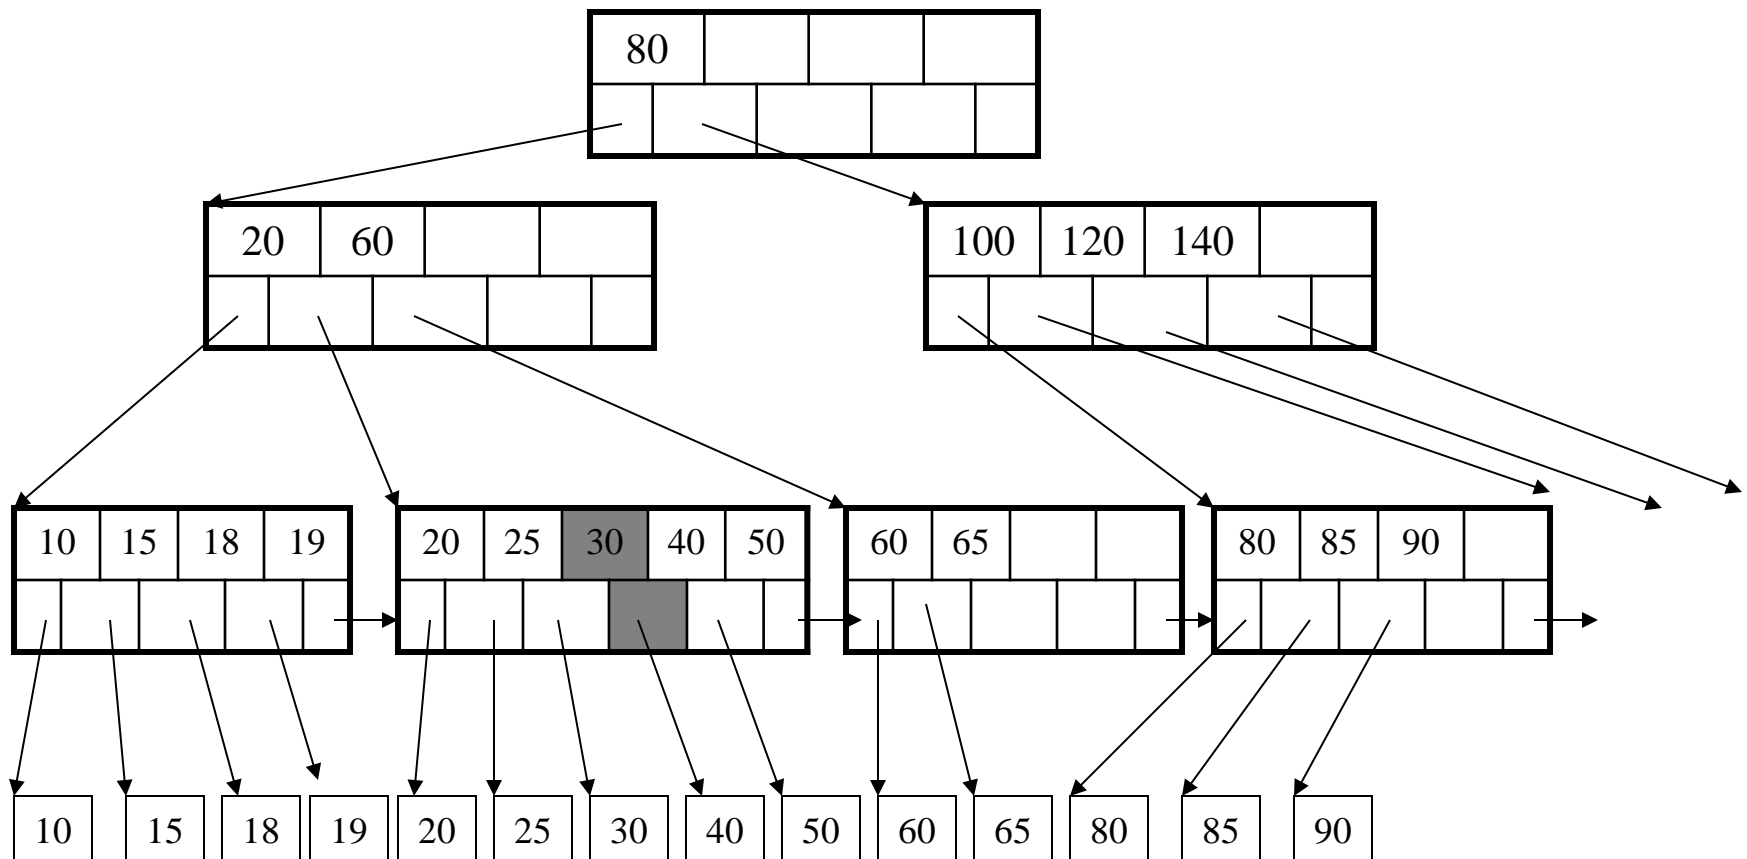
# Insertion into a B+ Tree

After insertion




# Insertion into a B+ Tree

But now have to split !



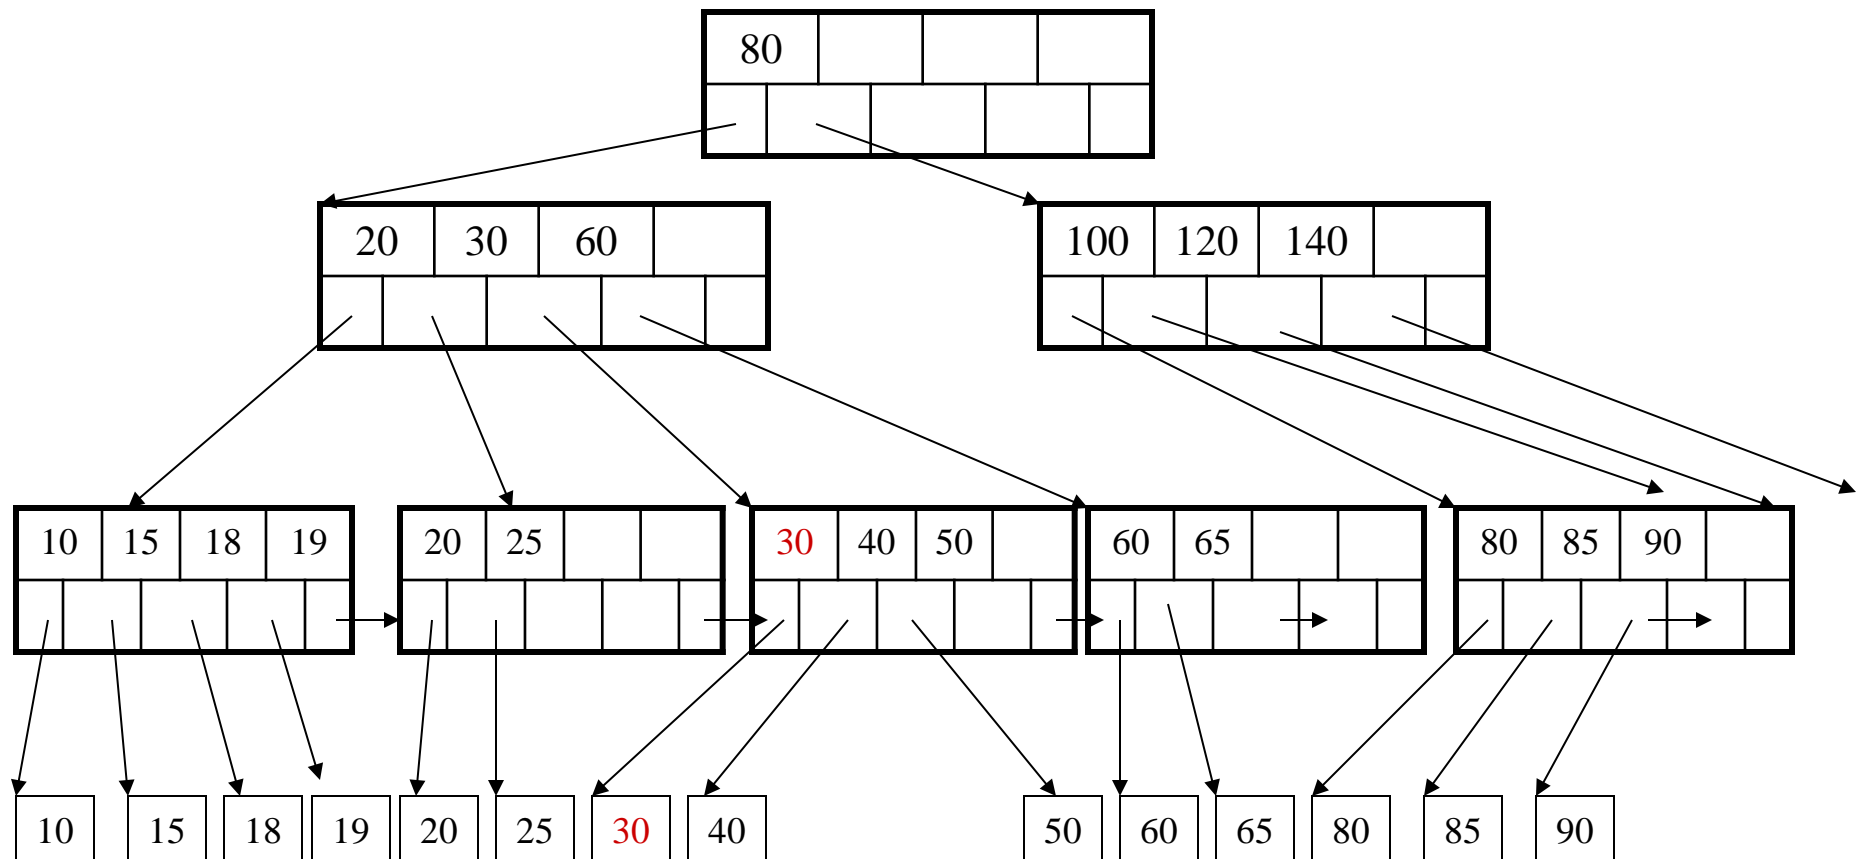
# Insertion into a B+ Tree

After the split



# Deletion from a B+ Tree


Delete 30



# Deletion from a B+ Tree

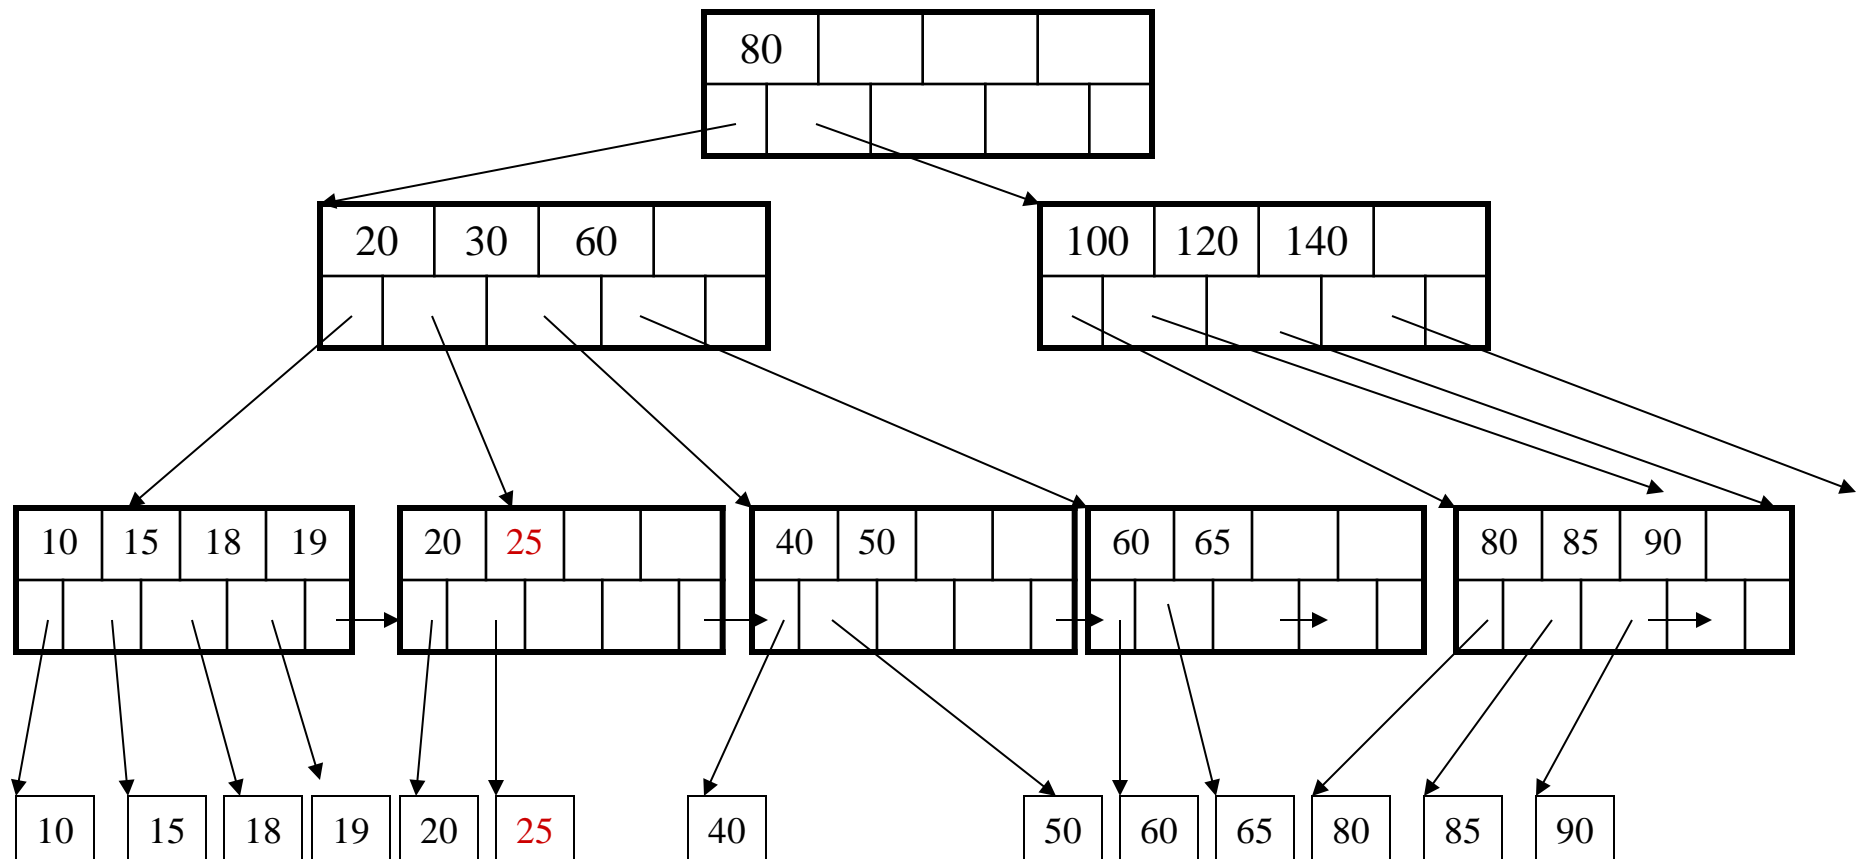
After deleting 30

May change to  
40, or not



# Deletion from a B+ Tree


Now delete 25



# Deletion from a B+ Tree

After deleting 25  
Need to rebalance

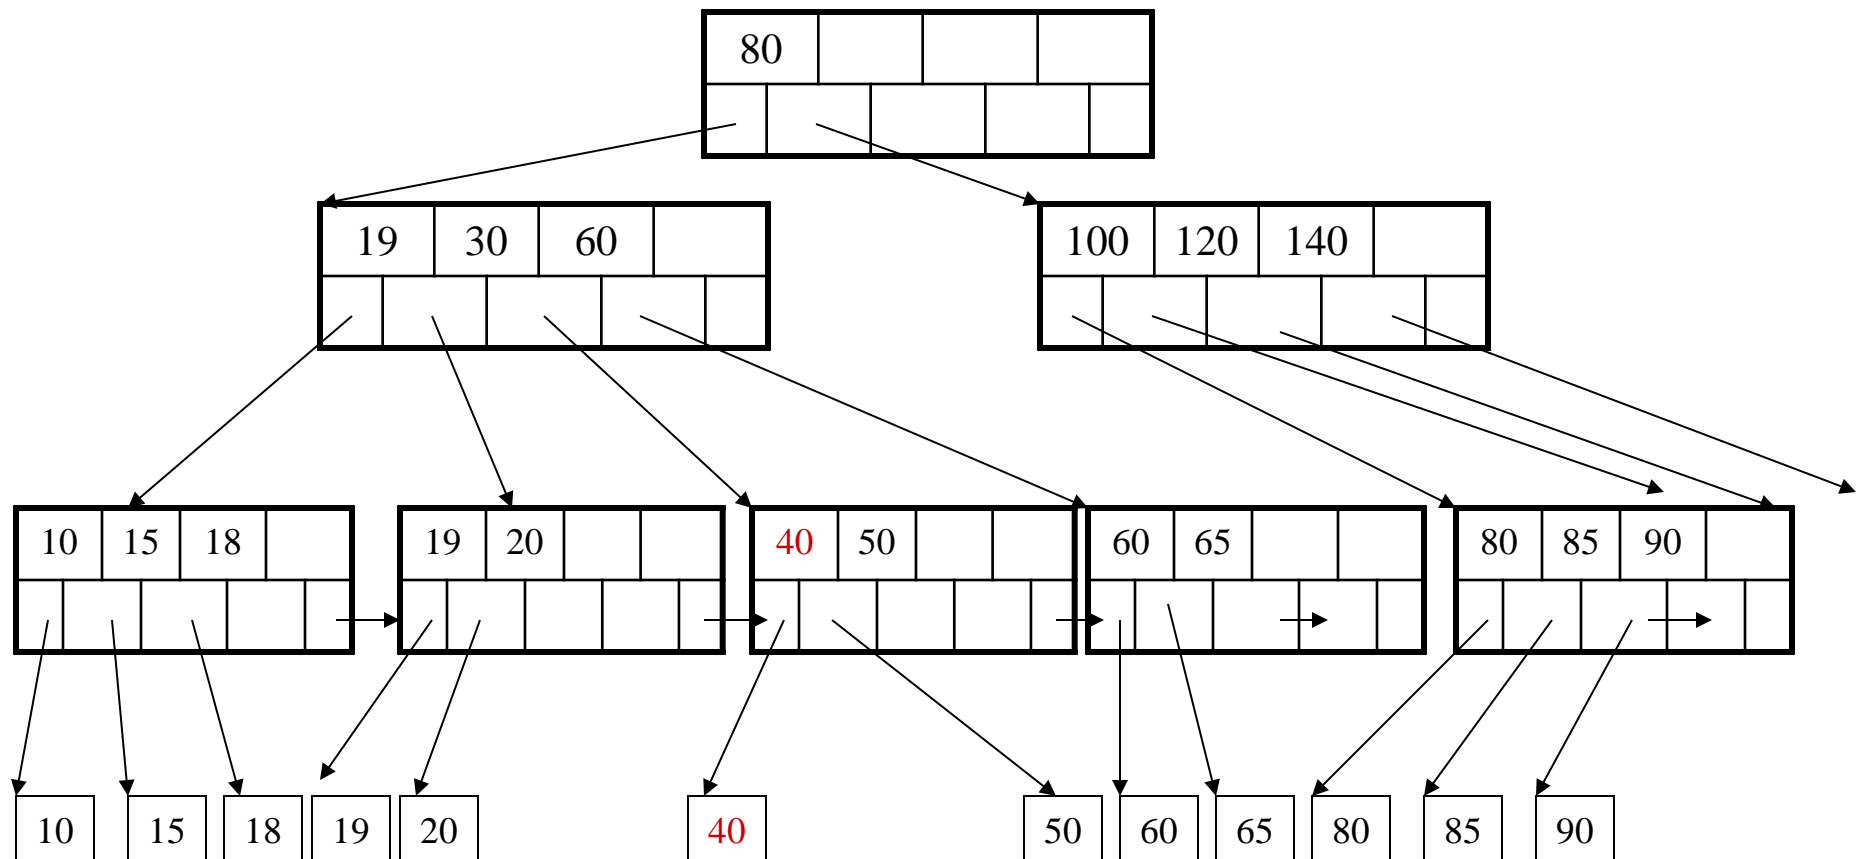
Rotate





# Deletion from a B+ Tree

Now delete 40




# Deletion from a B+ Tree

After deleting 40


Rotation not possible

Need to merge nodes



# Deletion from a B+ Tree

Final tree



# Deletion Strategy

- If a node is below the min capacity after deletion...
- Try the following in the given order
  1. move a key from **immediate** left sibling;
  2. move a key from immediate right sibling;
  3. merge with immediate left sibling;
  4. merge with immediate right sibling
- Cases 3 and 4 may lead to further removal of key from parent, and more fixing

# Another insertion example

