



INF 551

Wensheng Wu

Installation on EC2

- Create a new yum repository file for MongoDB
 - `cd /etc/yum.repos.d`
 - `sudo nano mongodb-org-3.4.repo`
- Add the following content to the file:
 - `[mongodb-org-3.4]`
 - `name=MongoDB Repository`
 - `baseurl=https://repo.mongodb.org/yum/amazon/2013.03/mongodb-org/3.4/x86_64/`
 - `gpgcheck=1`
 - `enabled=1`
 - `gpgkey=https://www.mongodb.org/static/pgp/server-3.4.asc`

Installation on EC2

- `sudo yum -y install mongodb-org`
- `sudo service mongod start`
 - Start the server
- `sudo service mongod stop`
 - Stop it

Document store

- MongoDB is a document database
- A document is similar to a JSON object
 - Consists of field-value pairs
 - Value may be another document, array, string, number, etc.
- Document = record/row in RDBMS

Collections

- Documents are stored in a collection
- Collection = table in RDBMS
- But documents may have different structures
 - In contrast, records in RDBMS have the same schema

Primary key

- Every document has a unique `_id` field
 - That acts as a primary key

MongoDB shell

- mongo

```
[ec2-user@ip-172-31-18-182 yum.repos.d]$ mongo
MongoDB shell version v3.4.9
connecting to: mongodb://127.0.0.1:27017
MongoDB server version: 3.4.9
Welcome to the MongoDB shell.
For interactive help, type "help".
For more comprehensive documentation, see
  http://docs.mongodb.org/
Questions? Try the support group
  http://groups.google.com/group/mongodb-user
Server has startup warnings:
2017-10-17T04:54:38.148+0000 I STORAGE [initandlisten]
2017-10-17T04:54:38.148+0000 I STORAGE [initandlisten] ** WARNING: Using
the XFS filesystem is strongly recommended with the wiredTiger storage
engine
2017-10-17T04:54:38.148+0000 I STORAGE [initandlisten] **           See
http://dochub.mongodb.org/core/prodnotes-filesystem
2017-10-17T04:54:38.225+0000 I CONTROL [initandlisten]
2017-10-17T04:54:38.225+0000 I CONTROL [initandlisten] ** WARNING: Access
control is not enabled for the database.
2017-10-17T04:54:38.225+0000 I CONTROL [initandlisten] **           Read
and write access to data and configuration is unrestricted.
2017-10-17T04:54:38.225+0000 I CONTROL [initandlisten]
> |
```

Create a new database

- No need to explicitly create it, just use it
 - It will be automatically created once you add a collection (i.e., table) to it

```
> show databases;
local 0.000GB
> use inf551
switched to db inf551
> show databases;
local 0.000GB
> use inf551
switched to db inf551
> db.createCollection('person')
{ "ok" : 1 }
> show databases;
inf551 0.000GB
local 0.000GB
```

```
> use inf551
switched to db inf551
> show collections
person
> show tables
person
> |
```


Databases

- use inf551
 - Switch to database "inf551"
- show databases
 - List all databases
- show tables/show collections
 - List all tables/collections in the current db
 - Can also say "show collections"

Database

- Dropping a database
 - `db.dropDatabase();`

Create/drop a collection

- `db.createCollection('person')`
 - `db` is a shell variable representing the current db
- `db.person.drop()`
 - Dropping a collection

Adding documents

- `db.person.insert({"_id": 1, "name": "john smith"})`
- `db.person.insert({"_id": 1, "name": "david smith"})`
 - Error: duplicate key!

ObjectId()

- ObjectId() function creates an ID
- `db.person.insert({"_id": ObjectId(), "name": "john smith"})`

```
writeResult({ "nInserted" : 1 })
> db.person.find()
{ "_id" : 1, "name" : "john smith" }
{ "_id" : ObjectId("58250aec7c61126eba98db48"), "name" : "john smith" }
```

ObjectId()

- `db.person.insert({"name": "john smith"})`
 - Here no specification of `"_id"` field
 - But an id will be automatically created

```
> db.person.find()
{ "_id" : 1, "name" : "john smith" }
{ "_id" : ObjectId("58250aec7c61126eba98db48"), "name" : "john smith" }
{ "_id" : ObjectId("58250d56249e740a9ddfbacc"), "name" : "john smith" }
> |
```

ObjectId()

- A 12-byte hexademical value
 - E.g., 58250aec7c61126eba98db48
- Among 12 bytes:
 - 4-byte: the seconds since 1970/1/1
 - 3-byte: machine identifier
 - 2-byte: process id
 - 3-byte: a counter, starting with a random value

Embedded sub-document

- `db.person.insert(
 {
 "name": "david johnson",
 "address": {"street": "123 maple",
 "city": "LA",
 "zip": 91989},
 "phone": ["323-123-0000", "626-124-0999"],
 "scores": [25, 35]
 })`

Array



Insert some more documents

- `db.person.insert({"name": "kevin small", "age": 35, "scores": [12, 20]})`
- `db.person.insert({"name": "mary lou", "age": 25})`

Query

- `db.person.find()`
 - Return all documents in person
- `db.person.find({"name": "kevin small"})`
 - Return all documents with specified name
- `db.person.find().pretty()`
 - Pretty print the output

Query operators

- Introduced by \$
- \$lt, \$gt, \$lte, \$gte, \$ne
 - Comparison operators
- \$or, \$and, \$not
 - Logical operators

Query operators

- `db.person.find({"age": {$gt: 25}})`
- `db.person.find({"name": "kevin small", "age": {$gt: 25}})`
 - Specify "and" condition
- `db.person.find({ $or: [{"name": "kevin small"}, {"age": {$gt: 25}}] })`
 - Specify "or" condition

Query operator

- What does each of these queries find?
 - `db.person.find({$or: [{"name":/kevin/i}, {"age":25}]})`
 - `db.person.find({$or: [{"name":/kevin/i, "age":25}]})`
 - `db.person.find({$and: [{"name":/kevin/i}, {"age":25}]})`

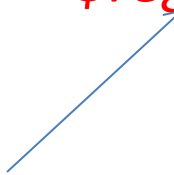
Query operator

- `db.person.find({name: {$not: {$eq: "john"}}})`

Same as:

- `db.person.find({name: {$ne: "john"}})`

Pattern matching

- `db.person.find({"name":/Kevin/i})`
 - This finds person whose name contains "kevin"
 - "i" means case-insensitive
 - Above is equivalent to:
 - `db.person.find({"name":{"$regex": /Kevin/, $options: 'i'}})`
 - In general, `/pattern/` where pattern is a regular expression
- \$regex is a query operator**
- 

Matching elements in array

- `db.person.find({"scores": {$gt: 20}})`
 - Note the "scores" field is an array and at least one value of the array should satisfy the specified condition (i.e., > 20)

Sorting

- `db.person.find().sort({age:-1})`
 - 1 for ascending; -1 descending
- Equivalent to:
 - Select *
 - From person
 - Order by age desc

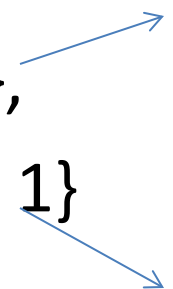
Limit

- `db.person.find().limit(1)`
 - Returns the first person

Distinct

- `db.person.distinct("age")`
- `db.person.distinct("age", {age: {$gt: 20}})`
 - distinct ages (for ages > 20)

Projection

- `db.person.find(`
 `{"age": {$ne: 25} },`
 `{"name":1, "age": 1}`
 `)`


Specify query condition

Specify projection
1: included in result; 0: do not
- This will return name and age (plus `_id`)
 - i.e., similar to 'select `_id`, name, age from users where age != 25'

Projection

- This does not work:
 - `db.person.find(
 {"age": {$ne: 25} },
 {"name":1, "age": 0}
)`
 - Can not mix 1 and 0 conditions (unless it is "_id")

Projection

- `db.person.find(
 {"age": {$ne: 25} },
 {"name":1, "age": 1, "_id": 0}
)`
- This does not return id, e.g.,
 `{ "name" : "john smith" }`
 `{ "name" : "david johnson" }`
 `{ "name" : "kevin small", "age" : 35 }`

Example

- Without projection

```
> db.person.find({"age": 25})
{ "_id" : ObjectId("582559b19f185cd8ccf23ff6"), "name" : "mary lou", "age" : 25 }
```

- With projection

```
> db.person.find({"age": 25}, {"name": 1, _id: 0})
{ "name" : "mary lou" }
```

Update documents

- `db.person.update(`
 `{ "age": { $gt: 25 } },`
 `{ $set: { "status": "C" } },`
 `{ multi: true }`
)
- Existing documents may not have status field; if not, insert it instead
- Update one or all documents


Similar to:

Update users set status = 'C' where age > 25

Another example

- `db.person.update({}, {$set: {"status":'C'}}, {multi:true})`
 - Note the empty query `{}`
- Add "status" field to all documents

Remove fields

- `db.person.update({}, {$unset: {"status": ""}}, {multi: true})`

Can put any value here
- Remove the "status" field from all documents

Remove documents

- `db.person.remove({})`
 - Remove all documents/records of person
- `db.person.remove({ "age": { $gt: 30 } })`
 - Remove documents which satisfy a condition

Remove a collection/table

- `db.person.drop()`
 - This will remove the person collection/table

Count()

- `db.person.count()`
 - Return # of documents in the person collection
- `db.person.count({age: {$gt: 25}})`
 - What does this do?
- `db.person.find({age: {$gt: 25}}).count()`

Query a embedded document

- Using **dot notation** to identify field in the embedded document
- `db.person.find({"address.city": "LA"})`
 - Return all documents whose city sub-field of address field = "LA"

Example for aggregation

- `db.product.insert({category: "cell", store:1, qty: 10})`
- `db.product.insert({category: "cell", store:2, qty: 20})`
- `db.product.insert({category: "laptop", store:1, qty: 10})`
- `db.product.insert({category: "laptop", store:2, qty: 30})`
- `db.product.insert({category: "laptop", store:2, qty: 40})`

Aggregation

- `db.product.aggregate([{$group: {_id: "$category", total:{$sum:"$qty"}}}])`
 - `{ "_id" : "laptop", "total" : 80 }`
 - `{ "_id" : "cell", "total" : 30 }`
- Similar to: `"select category, sum(qty) from product group by category"`

Aggregation

- `db.product.aggregate([{"$group": {"_id": "$category", "total":{"$sum:1"}} }])`
`{ "_id" : "laptop", "total" : 2 }`
`{ "_id" : "cell", "total" : 2 }`
- Similar to: "select category, count(*) from product group by category"

Aggregation with "having ..."

- `db.product.aggregate([{$group: {_id: "$category", total:{$sum:"$qty"}}}, {$match: {total: {$gt: 50}}}])`
 - `{ "_id" : "laptop", "total" : 80 }`
- In SQL:
Select category, sum(qty) total
from product
group by category
having total > 50

Aggregation on more than one field

- `db.product.aggregate([{$group: {_id: {cat: "$category", st: "$store"}, total: {$sum: "$qty"}}}])`

```
{ "_id" : { "cat" : "laptop", "st" : 1 }, "total" : 10 }  
{ "_id" : { "cat" : "laptop", "st" : 2 }, "total" : 70 }  
{ "_id" : { "cat" : "cell", "st" : 2 }, "total" : 20 }  
{ "_id" : { "cat" : "cell", "st" : 1 }, "total" : 10 }
```

Aggregation

- Other operators
 - \$avg
 - \$min
 - \$max

Aggregation pipeline

- `db.person.aggregate([{$match: {age: {$gt: 25}}}, {$group: {_id: "$gender", val: {$min: "$weight"}} }, {$match: {val: {$gt: 120}}}, {$limit: 2}, {$sort: {val: -1}}])`
- `$match -> $group -> $match -> $limit -> $sort`

Sharding in MongoDB

- Distribute documents/records in a large collection/table over multiple machines
- User can specify a sharding key
 - i.e., a field in a document
- Support sharding by key range or hashing

Sample data set

- Restaurants data
 - <https://raw.githubusercontent.com/mongodb/docs-assets/primer-dataset/primer-dataset.json>

Import sample dataset

- `mongoimport --db inf551 --collection restaurants --file primer-dataset.json`
 - No need to pre-create inf551 and restaurants if they do not exist yet
- More details:
 - <https://docs.mongodb.com/getting-started/shell/import-data/>

Resources

- Install MongoDB Community Edition on Amazon Linux
 - <https://docs.mongodb.com/manual/tutorial/install-mongodb-on-amazon/>